

Module 13: Using Window Ranking, Offset, and Aggregate Functions

Contents:

Module Overview

Lesson 1: Creating Windows with OVER

Lesson 2: Exploring Window Functions

Lab: Using Window Ranking, Offset, and Aggregate Functions

Module Review and Takeaways

Module Overview

Microsoft® SQL Server® implements support for SQL windowing operations, which means you can define a set of rows and apply several different functions against those rows. After you have learned how to work with windows and window functions, you might find that some types of queries that appeared to require complex manipulations of data (for example, self-joins, temporary tables, and other constructs) aren't needed to write your reports.

Objectives

After completing this module, you will be able to:

- Describe the benefits of using window functions.
- Restrict window functions to rows defined in an OVER clause, including partitions and frames.
- Write queries that use window functions to operate on a window of rows and return ranking, aggregation, and offset comparison results.

Lesson 1: Creating Windows with OVER

SQL Server provides a number of window functions, which perform calculations such as ranking, aggregations, and offset comparisons between rows. To use these functions, you will need to write queries that define windows, or sets, of rows. You will use the OVER clause and its related elements to define the sets for the window functions.

Lesson Objectives

After completing this lesson, you will be able to:

- Describe the T-SQL components used to define windows, and the relationships between them.
- Write queries that use the OVER clause, with partitioning, ordering, and framing to define windows.

SQL Windowing

- Windows extend T-SQL's set-based approach
- Windows allow you to specify an order as part of a calculation, without regard to order of input or final output order
- Windows allow partitioning and framing of rows to support functions
- Window functions can simplify queries that need to find running totals, moving averages, or gaps in data

```
SELECT Category, Qty, Orderyear,
       SUM(Qty) OVER (
         PARTITION BY category
         ORDER BY orderyear
         ROWS BETWEEN UNBOUNDED PRECEDING
         AND CURRENT ROW) AS RunningQty
FROM Sales.CategoryQtyYear;
```

SQL Server provides windows as a method for applying functions to sets of rows. There are many applications of this technique that solve common problems in writing T-SQL queries. For example, using windows allows the easy generation of row numbers in a result set and the calculation of running totals. Windows also provide an efficient way to compare values in one row with values in another without needing to join a table to itself using an inequality operator.

There are several core elements of writing queries that use windows:

1. Windows allow you to specify an order to rows that will be passed to a window function, without affecting the final order of the query output.
2. Windows include a partitioning feature, which enables you to specify that you want to restrict a function only to rows that have the same value as the current row.
3. Windows provide a framing option. It allows you to specify a further subset of rows within a window partition by setting upper and lower boundaries for the window frame, which presents rows to the window function.

Running Total Example

```
SELECT Category, Qty, Orderyear,
       SUM(Qty) OVER (PARTITION BY Category ORDER BY Orderyear
         ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW) AS RunningQty
FROM Sales.CategoryQtyYear;
```

The partial results:

Category	Qty	Orderyear	RunningQty
-----	-----	-----	-----
Beverages	1842	2006	1842
Beverages	3996	2007	5838
Beverages	3694	2008	9532
Condiments	962	2006	962
Condiments	2895	2007	3857
Condiments	1441	2008	5298
Confections	1357	2006	1357
Confections	4137	2007	5494
Confections	2412	2008	7906
Dairy Products	2086	2006	2086
Dairy Products	4374	2007	6460
Dairy Products	2689	2008	9149

During the next few topics of this lesson, you will learn how to use these query elements.


Windowing Components

- Conceptual relationship between window elements:

Result set
(OVER)

Window
partition
(PARTITION BY)

Frame
(ROWS BETWEEN)



In order to use windows and window functions in T-SQL, you will always use one of the subclauses that create and manipulate windows—the OVER subclause. Additionally, you may need to create partitions with the PARTITION BY option, and even further restrict which rows are applied to a function with framing options. Therefore, understanding the relationship between these components is vital.

The general relationship can be expressed as a sequence, with one element further manipulating the rows output by the previous element:

1. The OVER clause determines the result set that will be used by the window function. An OVER clause with no partition defined is unrestricted. It returns all rows to the function.
2. A PARTITION BY clause, if present, restricts the results to those rows with the same value in the partitioned columns as the current row. For example, PARTITION BY custid restricts the window to rows with the same custid as the current row. PARTITION BY builds on the OVER clause and cannot be used without OVER. (An OVER clause without a window partition clause is considered one partition).
3. A ROW or RANGE clause creates a window frame within the window partition, which allows you to set a starting and ending boundary around the rows being operated on. A frame requires an ORDER BY subclause within the OVER clause.

Windowing Example

```
SELECT Category, Qty, Orderyear,  
       SUM(Qty) OVER (PARTITION BY category ORDER BY Orderyear  
                     ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW) AS RunningQty  
FROM Sales.CategoryQtyYear;
```

The details of each component will be covered in future topics.

Note: A single query can use multiple window functions, each with its own OVER clause. Each clause determines its own partitioning, ordering, and framing.

Using OVER

- OVER defines a window, or set, of rows to be used by a window function, including any ordering
- With a specified window partition clause, the OVER clause restricts the set of rows to those with the same values in the partitioning elements
- By itself, OVER() is unrestricted and includes all rows
- Multiple OVER clauses can be used in a single query, each with its own partitioning and ordering, if needed

```
OVER ( [ <PARTITION BY clause> ]
      [ <ORDER BY clause> ]
      [ <ROWS or RANGE clause> ]
      )
```

The OVER clause defines the window, or set, of rows that will be operated on by a window function, which we will look at in the next lesson. The OVER clause includes partitioning, ordering, and framing, where each is applicable.

Used alone, the OVER clause does not restrict the result set passed to the window function. Used with a PARTITION BY subclause, OVER restricts the set to those rows with the same values in the partitioning elements.

OVER Example

```
SELECT Category, Qty, Orderyear,
       ROW_NUMBER() OVER (ORDER BY Qty DESC) AS Running
FROM Sales.CategoryQtyYear
ORDER BY Running;
```

The partial result, further ordered by the Running column for display purposes:

Category	Qty	Orderyear	Running
Dairy Products	4374	2007	1
Confections	4137	2007	2
Beverages	3996	2007	3
Beverages	3694	2008	4
Seafood	3679	2007	5
Condiments	2895	2007	6
Seafood	2716	2008	7

Dairy Products	2689	2008	8
Grains/Cereals	2636	2007	9

The next topics will build on this basic use of OVER to define a window of rows.

For further reading on the OVER clause, see Microsoft Docs:

OVER Clause (Transact-SQL)

<http://go.microsoft.com/fwlink/?LinkID=402778>

Partitioning Windows

- Partitioning limits a set to rows with the same value in the partitioning column
- Use PARTITION BY in the OVER() clause
- Without a PARTITION BY clause defined, OVER() creates a single partition of all rows

```
SELECT custid, ordermonth, qty,
       SUM(qty) OVER(PARTITION BY custid)
       AS totalbycust
FROM Sales.CustOrders;
```

custid	ordermonth	qty	totalbycust
1	2007-08-01 00:00:00.000	38	174
1	2007-10-01 00:00:00.000	41	174
2	2006-09-01 00:00:00.000	6	63
2	2007-08-01 00:00:00.000	18	63
3	2006-11-01 00:00:00.000	24	359
3	2007-04-01 00:00:00.000	30	359

Partitioning a window limits a set to rows with the same value in the partitioning column.

PARTITION BY Code Snippet

```
<function_name>() OVER(PARTITION BY Category)
```

As you have learned, if no partition is defined, then the OVER() clause returns all rows from the underlying query's result set to the window function.

PARTITION BY Example

```
SELECT Category, Qty, Orderyear,
       ROW_NUMBER() OVER (PARTITION BY Category ORDER BY Qty DESC) AS Running
FROM Sales.CategoryQtyYear
ORDER BY Category;
```

The partial result:

Category	Qty	Orderyear	Running
Beverages	3996	2007	1
Beverages	3694	2008	2
Beverages	1842	2006	3
Condiments	2895	2007	1
Condiments	1441	2008	2
Condiments	962	2006	3
Confections	4137	2007	1
Confections	2412	2008	2
Confections	1357	2006	3

Note: If you intend to add framing to the window partition, an ORDER BY subclause will also be needed in the OVER clause, as discussed in the next topic.

Ordering and Framing

- Window framing allows you to set start and end boundaries within a window partition
 - UNBOUNDED means go all the way to boundary in direction specified by PRECEDING or FOLLOWING (start or end)
 - CURRENT ROW indicates start or end at current row in partition
 - ROWS BETWEEN allows you to define a range of rows between two points
- Window ordering provides a context to the frame
 - Sorting by an attribute enables meaningful position of a boundary
 - Without ordering, "start at first row" is not useful because a set has no order

As you have learned, you use window partitions to define a subset of rows within the outer window defined by `OVER`. In a similar approach, window framing allows you to further restrict the rows available to the window function. You can think of a frame as a moving window over the data, starting and ending at positions you define.

To define window frames, use the `ROW` or `RANGE` subclauses to provide a starting and an ending boundary. For example, to set a frame that extends from the first row in the partition to the current row (such as to create a moving window for a running total), follow these steps:

1. Define an `OVER` clause with a `PARTITION BY` element.
2. Define an `ORDER BY` subclause to the `OVER` clause. This will cause the concept of "first row" to be meaningful.
3. Add the `ROWS BETWEEN` subclause, setting the starting boundary using `UNBOUNDED PRECEDING`. `UNBOUNDED` means go all the way to the boundary in the direction specified as `PRECEDING` (before). Add the `CURRENT ROW` element to indicate the ending boundary is the row being calculated.

Note: Since `OVER` returns a set, and sets have no order, an `ORDER BY` subclause is required for the framing operation to be useful. This can be (and typically is) different from `ORDER BY`, which determines the display order for the final result set.

Framing Example

```
SELECT Category, Qty, Orderyear,
       SUM(Qty) OVER (PARTITION BY Category ORDER BY Orderyear
                     ROWS BETWEEN UNBOUNDED PRECEDING
                     AND CURRENT ROW) AS RunningQty
FROM Sales.CategoryQtyYear;
```

The partial results:

Category	Qty	Orderyear	RunningQty
-----	-----	-----	-----
Beverages	1842	2006	1842
Beverages	3996	2007	5838
Beverages	3694	2008	9532
Condiments	962	2006	962
Condiments	2895	2007	3857
Condiments	1441	2008	5298
Confections	1357	2006	1357
Confections	4137	2007	5494
Confections	2412	2008	7906
Dairy Products	2086	2006	2086
Dairy Products	4374	2007	6460
Dairy Products	2689	2008	9149

Demonstration: Using OVER and Partitioning

In this demonstration, you will see how to use OVER, PARTITION BY, and ORDER BY clauses.

Demonstration Steps

Use OVER, PARTITION BY, and ORDER BY Clauses

1. Ensure that the **20761C-MIA-DC** and **20761C-MIA-SQL** virtual machines are both running, and then log on to **20761C-MIA-SQL** as **ADVENTUREWORKS\Student** with the password **Pa55w.rd**.
2. Run **D:\Demofiles\Mod13\Setup.cmd** as an administrator.
3. In the **User Account Control** dialog box, click **Yes**.
4. At the command prompt, type **y**, and then press Enter.
5. Wait for the script to finish, and then press any key.
6. Start SQL Server Management Studio and connect to the **MIA-SQL** database engine instance using Windows authentication.
7. Open the **Demo.ssmssln** solution in the **D:\Demofiles\Mod13\Demo** folder.
8. In Solution Explorer, open the **11 - Demonstration A.sql** script file.
9. Select the code under the comment **Step 1**, and then click **Execute**.
10. Select the code under the comment **Step 2**, and then click **Execute**.
11. Select the code under the comment **Step 3**, and then click **Execute**.
12. Select the code under the comment **Step 4**, and then click **Execute**.
13. Select the code under the comment **Step 5**, and then click **Execute**.
14. Keep SQL Server Management Studio open for the next demonstration.

Check Your Knowledge

Sequencing Activity

Put the following elements into the logical order in which they are processed in a windowed query by numbering each to indicate the correct order.

The OVER clause

The PARTITION BY clause

The ROW or RANGE clause

[Check answer](#)[Show solution](#)[Reset](#)**Correct**

Lesson 2: Exploring Window Functions

SQL Server provides window functions to operate on a window of rows. In addition to window aggregate functions, which you will find to be conceptually similar to grouped aggregate functions, you can use window ranking, distribution, and offset functions in your queries.

Lesson Objectives

After completing this lesson, you will be able to:

- Write queries that use window aggregate functions.
- Write queries that use window ranking functions.
- Write queries that use window offset functions.

Defining Window Functions

- A window function is a function applied to a window, or set, of rows
- Window functions include aggregate, ranking, distribution, and offset functions
- Window functions depend on set created by `OVER()`

```
SELECT productid, productname, unitprice,  
       RANK() OVER(ORDER BY unitprice DESC)  
       AS pricerank  
FROM Production.Products  
ORDER BY pricerank;
```

A window function is applied to a window, or set, of rows. Earlier in this course, you learned about group aggregate functions such as `SUM`, `MIN`, and `MAX`, which operated on a set of rows defined by a `GROUP BY` clause. In

window operations, you can use these functions, in addition to others, to operate on a set of rows defined in a window by an OVER clause and its elements.

SQL Server window functions can be found in the following categories, which will be discussed in the next topics:

- Aggregate functions, such as SUM, which operate on a window and return a single row.
- Ranking functions, such as RANK, which depend on a sort order and return a value representing the rank of a row, with respect to other rows in the window.
- Distribution functions, such as CUME_DIST, which calculate the distribution of a value in a window of rows.
- Offset functions, such as LEAD, which return values from other rows relative to the position of the current row.

When used in windowing scenarios, these functions depend on the result set returned by the OVER clause and any further restrictions you provide within OVER, such as partitioning and framing.

RANK Example

```
SELECT productid, productname, unitprice,
       RANK() OVER(ORDER BY unitprice DESC) AS pricerank
FROM Production.Products
ORDER BY pricerank;
```

The partial result:

productid	productname	unitprice	pricerank
38	Product QDOMO	263.50	1
29	Product VJXYN	123.79	2
9	Product AOZBW	97.00	3
20	Product QHFFP	81.00	4
18	Product CKEDC	62.50	5
59	Product UKXRI	55.00	6
51	Product APITJ	53.00	7
62	Product WUXYK	49.30	8
43	Product ZZZHR	46.00	9
28	Product OFBNT	45.60	10
27	Product SMIOH	43.90	11
63	Product ICKNK	43.90	11
8	Product WVJFP	40.00	13

RANK with PARTITION Example

```

SELECT categoryid, productid, unitprice,
       RANK() OVER(PARTITION BY categoryid ORDER BY unitprice DESC) AS pricerank
FROM Production.Products
ORDER BY categoryid, pricerank;

```

The partial result, edited for space:

categoryid	productid	unitprice	pricerank
1	38	263.50	1
1	43	46.00	2
1	2	19.00	3
2	63	43.90	1
2	8	40.00	2
2	61	28.50	3
2	6	25.00	4
3	20	81.00	1
3	62	49.30	2
3	27	43.90	3
3	26	31.23	4

Notice that the addition of partitioning allows the window function to operate at a more granular level than when OVER returns an unrestricted set.

Note: Repeating values and gaps in the pricerank column are expected when using RANK in case of ties. Use DENSE_RANK if gaps are not desired. See the next topics for more information.

Window Aggregate Functions

- Similar to grouped aggregate functions
 - SUM, MIN, MAX, and so on
- Applied to windows defined by OVER clause
- Window aggregate functions support partitioning, ordering, and framing

```
SELECT custid, ordermonth, qty,
       SUM(qty) OVER(PARTITION BY custid)
       AS totalpercust
FROM Sales.CustOrders;
```

Window aggregate functions are similar to the aggregate functions you have already used in this course. They aggregate a set of rows and return a single value. However, when used in the context of windows, they operate on the set returned by the OVER clause, not on a set defined by a grouped query using GROUP BY.

Window aggregate functions provide support for windowing elements you have learned about in this module, such as partitioning, ordering, and framing. Unlike other window functions, ordering is not required with aggregate functions, unless you are also specifying a frame.

Window Aggregate Example

```
SELECT custid,
       ordermonth,
       qty,
       SUM(qty) OVER ( PARTITION BY custid ) AS totalpercust
FROM Sales.CustOrders;
```

The partial result, edited for space:

custid	ordermonth	qty	totalpercust
1	2007-08-01 00:00:00.000	38	174
1	2007-10-01 00:00:00.000	41	174
1	2008-01-01 00:00:00.000	17	174
2	2006-09-01 00:00:00.000	6	63
2	2007-08-01 00:00:00.000	18	63

3	2006-11-01 00:00:00.000	24	359
3	2007-04-01 00:00:00.000	30	359
3	2007-05-01 00:00:00.000	80	359
4	2007-02-01 00:00:00.000	40	650
4	2007-06-01 00:00:00.000	96	650

Further Window Aggregate Example

```

SELECT  custid, ordermonth, qty,
        SUM(qty) OVER ( PARTITION BY custid ) AS custtotal,
        CAST(100. * qty/SUM(qty) OVER ( PARTITION BY custid )AS NUMERIC(8,2)) AS
ofTotal
FROM    Sales.CustOrders;

```

The result:

custid	ordermonth	qty	custtotal	ofTotal
1	2007-08-01 00:00:00.000	38	174	21.84
1	2007-10-01 00:00:00.000	41	174	23.56
1	2008-01-01 00:00:00.000	17	174	9.77
1	2008-03-01 00:00:00.000	18	174	10.34
1	2008-04-01 00:00:00.000	60	174	34.48
2	2006-09-01 00:00:00.000	6	63	9.52
2	2007-08-01 00:00:00.000	18	63	28.57
2	2007-11-01 00:00:00.000	10	63	15.87
2	2008-03-01 00:00:00.000	29	63	46.03
3	2006-11-01 00:00:00.000	24	359	6.69
3	2007-04-01 00:00:00.000	30	359	8.36
3	2007-05-01 00:00:00.000	80	359	22.28
3	2007-06-01 00:00:00.000	83	359	23.12
3	2007-09-01 00:00:00.000	102	359	28.41
3	2008-01-01 00:00:00.000	40	359	11.14

Window Ranking Functions

- Ranking functions require a window order clause
- Partitioning is optional
- To display results in sorted order still requires ORDER BY!

Function	Description
RANK	Returns the rank of each row within the partition of a result set. May include ties and gaps.
DENSE_RANK	Returns the rank of each row within the partition of a result set. May include ties. Will not include gaps.
ROW_NUMBER	Returns a unique sequential row number within partition based on current order.
NTILE	Distributes the rows in an ordered partition into a specified number of groups. Returns the number of the group to which the current row belongs.

Window ranking functions return a value representing the rank of a row with respect to other rows in the window. To accomplish this, ranking functions require an ORDER BY element within the OVER clause, to establish the position of each row within the window.

Note: Remember that the ORDER BY element within the OVER clause affects only the processing of rows by the window function. To control the display order of the results, add an ORDER BY clause to the end of the SELECT statement, as with other queries.

The primary difference between RANK and DENSE_RANK is the handling of rows when there are tie values.

RANK and DENSE_RANK Example

```
SELECT CatID, CatName, ProdName, UnitPrice,
       RANK() OVER(PARTITION BY CatID ORDER BY UnitPrice DESC) AS PriceRank,
       DENSE_RANK() OVER(PARTITION BY CatID ORDER BY UnitPrice DESC) AS DensePriceRank
FROM Production.CategorizedProducts
ORDER BY CatID;
```

The partial results follow. Note the rank numbering of the rows following the products with a unitprice of 18.00:

CatID	CatName	ProdName	UnitPrice	PriceRank	DensePriceRank
1	Beverages	Product QDOMO	263.50	1	1
1	Beverages	Product ZZZHR	46.00	2	2

1	Beverages	Product	RECZE	19.00	3	3
1	Beverages	Product	HHYDP	18.00	4	4
1	Beverages	Product	LSOFL	18.00	4	4
1	Beverages	Product	NEVTJ	18.00	4	4
1	Beverages	Product	JYGFE	18.00	4	4
1	Beverages	Product	TOONT	15.00	8	5
1	Beverages	Product	XLXQF	14.00	9	6
1	Beverages	Product	SWNJY	14.00	9	6
1	Beverages	Product	BWRLG	7.75	11	7
1	Beverages	Product	QOGNU	4.50	12	8

Go to Ranking Functions (Transact-SQL) in the SQL Server Technical Documentation:

Ranking Functions (Transact-SQL)

<http://go.microsoft.com/fwlink/?LinkID=402779>

Window Distribution Functions

- Window distribution functions perform statistical analysis on data, and require a window order clause
- Rank distribution performed with PERCENT_RANK and CUME_DIST
- Inverse distribution performed with PERCENTILE_CONT and PERCENTILE_DISC

Window distribution functions perform statistical analysis on the rows within the window or window partition. Partitioning a window is optional for distribution functions, but ordering is required.

Distribution functions return a rank of a row, but instead of being expressed as an ordinal number, as with RANK, DENSE_RANK, or ROW_NUMBER, it is expressed as a ratio between 0 and 1. SQL Server provides rank distribution with the PERCENT_RANK and CUME_DIST functions. It provides inverse distribution with the PERCENTILE_CONT and PERCENTILE_DISC functions.

These functions are listed here for completeness only and are beyond the scope of this course. For more information, see the SQL Server Technical Documentation:

Analytic Functions (Transact-SQL)

<http://go.microsoft.com/fwlink/?LinkID=402780>

Window Offset Functions

- Window offset functions allow comparisons between rows in a set without the need for a self-join
- Offset functions operate on a position relative to the current row, or to the start or end of the window frame

Function	Description
LAG	Returns an expression from a previous row that is a defined offset from the current row. Returns NULL if no row at specified position.
LEAD	Returns an expression from a later row that is a defined offset from the current row. Returns NULL if no row at specified position.
FIRST_VALUE	Returns the first value in the current window frame. Requires window ordering to be meaningful.
LAST_VALUE	Returns the last value in the current window frame. Requires window ordering to be meaningful.

Windows offset functions give access to values located in rows other than the current row. This can enable queries that perform comparisons between rows, without the need to join the table to itself.

Offset functions operate on a position that is either relative to the current row, or relative to the starting or ending boundary of the window frame. LAG and LEAD operate on an offset to the current row. FIRST_VALUE and LAST_VALUE operate on an offset from the window frame.

Note: Since FIRST_VALUE and LAST_VALUE operate on offsets from the window frame, it is important to remember to specify framing options other than the default of RANGE BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW.

Window Offset Function Example

```
SELECT employee, orderyear ,totalsales AS currsales,
       LEAD(totalsales, 1,0) OVER (PARTITION BY employee ORDER BY orderyear) AS nextsales
FROM Sales.OrdersByEmployeeYear
ORDER BY employee, orderyear;
```

The partial results:

employee	orderyear	currsales	nextsales
1	2006	38789.00	97533.58
1	2007	97533.58	65821.13
1	2008	65821.13	0.00
2	2006	22834.70	74958.60
2	2007	74958.60	79955.96
2	2008	79955.96	0.00
3	2006	19231.80	111788.61
3	2007	111788.61	82030.89
3	2008	82030.89	0.00

Demonstration: Exploring Windows Functions

In this demonstration, you will see how to use window aggregate, ranking, and offset functions.

Demonstration Steps

Use Window Aggregate, Ranking, and Offset Functions

1. In Solution Explorer, open the **21 - Demonstration B.sql** script file.
2. Select the code under the comment **Step 1**, and then click **Execute**.
3. Select the code under the comment **Step 2**, and then click **Execute**.
4. Select the code under the comment **Step 3**, and then click **Execute**.
5. Select the code under the comment **Step 4**, and then click **Execute**.
6. Select the code under the comment **Step 5**, and then click **Execute**.
7. Select the code under the comment **Step 6**, and then click **Execute**.
8. Select the code under the comment **Step 7**, and then click **Execute**.
9. Select the code under the comment **Step 8**, and then click **Execute**.
10. Select the code under the comment **Step 9**, and then click **Execute**.
11. Select the code under the comment **Step 10**, and then click **Execute**.
12. Select the code under the comment **Step 11**, and then click **Execute**.
13. Close SQL Server Management Studio without saving any files.

Check Your Knowledge

Categorize Activity

Place each windowing function into the appropriate category. Indicate your answer by writing the category number to the right of each item.

Window Aggregate Functions

MIN()

MAX()

SUM()

Window Ranking Function

DENSERANK()

ROW_NUMBER()

NTITLE()

RANK()

Window Distribution Functions

PERCENTILE_CONT()

CUME_DIST()

PERCENTILE_DISC()

PERCENT_RANK()

Check answer

Show solution

Reset

Correct

Lab: Using Window Ranking, Offset, and Aggregate Functions

Scenario

As a business analyst for Adventure Works, you will be writing reports using corporate databases stored in SQL Server. You have been provided with a set of business requirements for data and you will write T-SQL queries to retrieve the specified data from the databases. To fill these requests, you will need to calculate ranking values, as well as the difference between two consecutive rows, and running totals. You will use window functions to achieve these calculations.

Objectives

After completing this lab, you will be able to:

- Write queries that use ranking functions.
- Write queries that use offset functions.
- Write queries that use window aggregation functions.

Lab Setup

Estimated Time: 60 minutes

Virtual machine: **20761C-MIA-SQL**

User name: **ADVENTUREWORKS\Student**

Password: **Pa55w.rd**

Exercise 1: Writing Queries That Use Ranking Functions

Scenario

The sales department would like to rank orders by their values for each customer. You will provide the report by using the RANK function. You will also practice how to add a calculated column to display the row number in the SELECT clause.

The main tasks for this exercise are as follows:

1. Prepare the Lab Environment
2. Write a SELECT Statement That Uses the ROW_NUMBER Function to Create a Calculated Column
3. Add an Additional Column Using the RANK Function
4. Write A SELECT Statement to Calculate a Rank, Partitioned by Customer and Ordered by the Order Value
5. Write a SELECT Statement to Rank Orders, Partitioned by Customer and Order Year, and Ordered by the Order Value

6. Filter Only Orders with the Top Two Ranks



Detailed Steps ▲

Task 1: Prepare the Lab Environment

1. Ensure that the **20761C-MIA-DC** and **20761C-MIA-SQL** virtual machines are both running, and then log on to **20761C-MIA-SQL** as **ADVENTUREWORKS\Student** with the password **Pa55w.rd**.
2. Run **Setup.cmd** in the **D:\Labfiles\Lab13\Starter** folder as Administrator.



Detailed Steps ▲

Task 2: Write a SELECT Statement That Uses the ROW_NUMBER Function to Create a Calculated Column

1. In SQL Server Management Studio, open the project file **D:\Labfiles\Lab13\Starter\Project\Project.ssmssln** and the T-SQL script **51 - Lab Exercise 1.sql**. Ensure that you are connected to the TSQL database.
2. Write a SELECT statement to retrieve the orderid, orderdate, and val columns in addition to a calculated column named rowno from the view Sales.OrderValues. Use the ROW_NUMBER function to return rowno. Order the row numbers by the orderdate column.
3. Execute the written statement and compare the results that you achieved with the desired results shown in the file D:\Labfiles\Lab13\Solution\52 - Lab Exercise 1 - Task 1 Result.txt.



Detailed Steps ▲

Task 3: Add an Additional Column Using the RANK Function

1. Copy the previous T-SQL statement and modify it by including an additional column named rankno. To create rankno, use the RANK function, with the rank order based on the orderdate column.
2. Execute the modified statement and compare the results that you achieved with the desired results shown in the file D:\Labfiles\Lab13\Solution\53 - Lab Exercise 1 - Task 2 Result.txt. Notice the different values in the rowno and rankno columns for some of the rows.
3. What is the difference between the RANK and ROW_NUMBER functions?



Detailed Steps ▲

Task 4: Write A SELECT Statement to Calculate a Rank, Partitioned by Customer and Ordered by the Order Value

1. Write a SELECT statement to retrieve the orderid, orderdate, custid, and val columns, as well as a calculated column named orderrankno from the Sales.OrderValues view. The orderrankno column should display the rank per each customer independently, based on val ordering in descending order.
2. Execute the written statement and compare the results that you achieved with the desired results shown in the file D:\Labfiles\Lab13\ Solution\54 - Lab Exercise 1 - Task 3 Result.txt.



Detailed Steps ▲

Task 5: Write a SELECT Statement to Rank Orders, Partitioned by Customer and Order Year, and Ordered by the Order Value

1. Write a SELECT statement to retrieve the custid and val columns from the Sales.OrderValues view. Add two calculated columns:
 - o orderyear as a year of the orderdate column.
 - o orderrankno as a rank number, partitioned by the customer and order year, and ordered by the order value in descending order.
2. Execute the written statement and compare the results that you achieved with the desired results shown in the file D:\Labfiles\Lab13\ Solution\55 - Lab Exercise 1 - Task 4 Result.txt.



Detailed Steps ▲

Task 6: Filter Only Orders with the Top Two Ranks

1. Copy the previous query and modify it to filter only orders with the first two ranks based on the orderrankno column.
2. Execute the written statement and compare the results that you achieved with the desired results shown in the file D:\Labfiles\Lab13\ Solution\56 - Lab Exercise 1 - Task 5 Result.txt.

Result: After this exercise, you should know how to use ranking functions in T-SQL statements.

Exercise 2: Writing Queries That Use Offset Functions

Scenario

You need to provide separate reports to analyze the difference between two consecutive rows. This will enable business users to analyze growth and trends.

The main tasks for this exercise are as follows:

1. Write a SELECT Statement to Retrieve the Next Row Using a Common Table Expression (CTE)
2. Add a Column to Display the Running Sales Total
3. Analyze the Sales Information for the Year 2007



Detailed Steps ▲

Task 1: Write a SELECT Statement to Retrieve the Next Row Using a Common Table Expression (CTE)

1. Open the T-SQL script **71 - Lab Exercise 3.sql**. Ensure that you are connected to the TSQL database.
2. Define a CTE named OrderRows based on a query that retrieves the orderid, orderdate, and val columns from the Sales.OrderValues view. Add a calculated column named rowno using the ROW_NUMBER function, ordering by the orderdate and orderid columns.
3. Write a SELECT statement against the CTE and use the LEFT JOIN with the same CTE to retrieve the current row and the previous row based on the rowno column. Return the orderid, orderdate, and val columns for the current row and the val column from the previous row as prevval. Add a calculated column named diffprev to show the difference between the current val and previous val.
4. Execute the T-SQL code and compare the results that you achieved with the desired results shown in the file D:\Labfiles\Lab13\Solution\62 - Lab Exercise 2 - Task 1 Result.txt.



Detailed Steps ▲

Task 2: Add a Column to Display the Running Sales Total

1. Write a SELECT statement that uses the LAG function to achieve the same results as the query in the previous task. The query should not define a CTE.
2. Execute the written statement and compare the results that you achieved with the recommended results shown in the file D:\Labfiles\Lab13\Solution\63 - Lab Exercise 2 - Task 2 Result.txt.



Detailed Steps ▲

Task 3: Analyze the Sales Information for the Year 2007

1. Define a CTE named SalesMonth2007 that creates two columns: monthno (the month number of the orderdate column) and val (aggregated val column). Filter the results to include only the order year 2007 and group by monthno.
2. Write a SELECT statement to retrieve the monthno and val columns. Add two calculated columns:
 - o **avglast3months**. This column should contain the average sales amount for the last three months before the current month, using a window aggregate function. You can assume that there are no missing months.
 - o **ytdval**. This column should contain the cumulative sales value up to the current month.
3. Execute the written statement and compare the results that you achieved with the recommended results shown in the file D:\Labfiles\Lab13\ Solution\63 - Lab Exercise 2 - Task 3 Result.txt.

Result: After this exercise, you should be able to use the offset functions in your T-SQL statements.

Exercise 3: Writing Queries That Use Window Aggregate Functions

Scenario

To better understand the cumulative sales value of a customer through time and to provide the sales analyst with a year-to-date analysis, you will have to write different SELECT statements that use the window aggregate functions.

The main tasks for this exercise are as follows:

1. Write a SELECT Statement to Display the Contribution of Each Customer's Order Compared to That Customer's Total Purchase
2. Add a Column to Display the Running Sales Total
3. Analyze the Year-to-Date Sales Amount and Average Sales Amount for the Last Three Months



Detailed Steps ▲

Task 1: Write a SELECT Statement to Display the Contribution of Each Customer's Order Compared to That Customer's Total Purchase

1. Open the T-SQL script **71 - Lab Exercise 3.sql**. Ensure that you are connected to the TSQL database.
2. Write a SELECT statement to retrieve the custid, orderid, orderdate, and val columns from the Sales.OrderValues view. Add a calculated column named percoftotalcust containing a percentage value of each order sales amount compared to the total sales amount for that customer.

3. Execute the written statement and compare the results that you achieved with the recommended results shown in the file D:\Labfiles\Lab13\ Solution\72 - Lab Exercise 3 - Task 1 Result.txt.



Detailed Steps ▲

Task 2: Add a Column to Display the Running Sales Total

1. Copy the previous SELECT statement and modify it by adding a new calculated column named `runval`. This column should contain a running sales total for each customer based on order date, using `orderid` as the tiebreaker.
2. Execute the written statement and compare the results that you achieved with the recommended results shown in the file D:\Labfiles\Lab13\ Solution\73 - Lab Exercise 3 - Task 2 Result.txt.



Detailed Steps ▲

Task 3: Analyze the Year-to-Date Sales Amount and Average Sales Amount for the Last Three Months

1. Copy the `SalesMonth2007` CTE in the final task in exercise 2. Write a SELECT statement to retrieve the `monthno` and `val` columns. Add two calculated columns:
 - o **avglast3months**. This column should contain the average sales amount for the last three months before the current month using a window aggregate function. You can assume that there are no missing months.
 - o **ytdval**. This column should contain the cumulative sales value up to the current month.
2. Execute the written statement and compare the results that you achieved with the recommended results shown in the file D:\Labfiles\Lab13\ Solution\74 - Lab Exercise 3 - Task 3 Result.txt.
3. Close SQL Server Management Studio without saving any changes.

Result: After this exercise, you should have a basic understanding of how to use window aggregate functions in T-SQL statements.

Module Review and Takeaways

In this module, you have learned how to:

- Describe the benefits of using window functions.

- Restrict window functions to rows defined in an OVER clause, including partitions and frames.
- Write queries that use window functions to operate on a window of rows and return ranking, aggregation, and offset comparison results.

Review Question(s)

Check Your Knowledge

Discovery

What results will be returned by a ROW_NUMBER function if there is no ORDER BY clause in the query?

Show solution

Reset

An unordered set.

Check Your Knowledge

Discovery

Which ranking function would you use to return the values 1,1,3? Which would return 1,1,2?

Show solution

Reset

RANK, DENSE_RANK.

Check Your Knowledge

Discovery

Can a window frame extend beyond the boundaries of the window partition defined in the same OVER() clause?

Show solution

Reset

No.